
ShadowWeave: Calibrated Amodal Occupancy Completion into Occluded Space from a Single Depth Frame

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A mobile agent must reason about the space occluded by the very obstacles it
2 navigates around (“shadow”); we study how well a model can **complete** occupancy
3 there from a **single monocular-depth frame**, and whether it knows *when it is*
4 *guessing*. ShadowWeave projects one depth frame into an egocentric bird’s-eye-
5 view (BEV) occupancy grid, marks unobserved cells with an explicit shadow mask,
6 and predicts occupancy and per-cell uncertainty inside that mask. On a MuJoCo
7 indoor benchmark, the model beats a persistence baseline by a micro-averaged
8 shadow-IoU gain of **+0.31 to +0.36 at a 5 s horizon** (95% per-episode bootstrap
9 CIs clear of zero at every horizon), far exceeds an observation-blind dataset prior,
10 and its uncertainty is **better calibrated inside shadow (ECE 0.043–0.047) than**
11 **in directly observed space (0.058–0.086)**. A decomposition separating *amodal*
12 *completion* from *temporal forecasting* shows the gain is almost entirely completion;
13 the pure-forecasting residual sits at the noise floor. Randomizing room geometry
14 per episode leaves the gain statistically unchanged on never-seen rooms, ruling
15 out memorization. The completed grid and its uncertainty drive an eyes-free
16 spatial-audio interface at 7 ms (p95) perception-to-planning latency; we release the
17 methodology as a reusable protocol for honest occupancy-completion claims.

18 1 Introduction

19 Agents — a person entering a corridor, a robot rounding a shelf, a blind traveler crossing a room
20 — act in space they cannot see; for safety-critical, eyes-free use, inference about that space needs
21 an honest estimate of its own reliability: a confident hallucination of clear space behind a wall is
22 worse than a calibrated “I don’t know.” Prior work addresses parts of this problem but not their
23 conjunction (§5). **ShadowWeave** produces, from a single monocular-depth frame, (i) an egocentric
24 BEV occupancy grid, (ii) an explicit shadow mask marking cells the sensor cannot see, and (iii) a
25 calibrated per-cell occupancy probability inside that mask. Our contributions are:

- 26 • **A calibrated amodal-completion result.** From one depth frame, the model completes
27 occluded occupancy with a micro-averaged shadow gain of +0.31–0.36 at 5 s (CIs clear
28 of zero at every horizon), and its in-shadow uncertainty is *better* calibrated (ECE 0.043–
29 0.047) than in observed space — the model conveys, not merely produces, hidden-space
30 predictions.
- 31 • **A decomposition methodology** that separates completion from forecasting. It reveals that
32 the gain is spatial completion of persistent hidden structure; the pure-forecasting residual




Figure 1: Pipeline: depth frame $\rightarrow$ 96×96 egocentric BEV occupancy and visibility; thresholded visibility gives the shadow mask; a 31M U-Net predicts occupancy at 1/3/5/10 s inside it, feeding A* planning and nine-zone HRTF audio.

is at the noise floor (fixed static +0.000; moving tier +0.007 [+0.002, +0.012] at 5 s — significant but practically small, decaying with horizon as real forecasting should). We report this negative honestly: most pooled occupancy numbers in the literature share exactly this confound.

- **A memorization confound-kill.** Randomizing room geometry per episode leaves the gain statistically unchanged on never-seen rooms (+0.306 fixed vs +0.309 [+0.294, +0.324] randomized), showing completion is per-scene inference, not a memorized layout.
- **A real-time eyes-free system** that turns completed occupancy and its uncertainty into an A* path and nine-zone spatial-audio cues at 7 ms p95 perception-to-planning latency.

2 Method

A differentiable projector (Fig. 1) lifts each monocular-depth image (rendered in simulation; from a monocular-depth network in deployment) into an egocentric 96×96 BEV occupancy grid covering 5 m ahead, plus a **visibility** grid: cells along a ray up to the first returned surface are observed-free, the surface cell observed-occupied, and everything beyond is **shadow**. Thresholding visibility at 0.5 yields the shadow mask — the region of interest for every metric here; optical flow between BEV frames is an optional channel. A 31M-parameter U-Net predicts occupancy from the single-frame BEV stack at 1/3/5/10 s as per-cell probabilities, trained with positive-up-weighted binary cross-entropy (BEV targets are $\approx 7\%$ occupied) on 400 episodes with early stopping and EMA weights; we report the epoch-8 checkpoint (val. loss 0.300, IoU@5s 0.665). The input is a **single frame** with no multi-frame memory, making the completion-vs-forecasting distinction (§3) sharp. The per-cell probability *is* the uncertainty estimate; we test calibration **inside the shadow region specifically**, not merely grid-wide (where easy free space dominates). Downstream, the completed grid and its uncertainty drive an A* planner (extra cost through unobserved cells), a reactive collision-avoidance policy, and nine-zone HRTF audio cues; the audio interface is a demonstration, not user-evaluated.

3 Evaluation methodology

The empty-credit trap and micro-averaging. Our claim lives entirely inside a sparse, mostly-empty sub-region (shadow is $\sim 92\%$ never-occupied), and the natural masked IoU scores an empty prediction against an empty masked region as a perfect 1.0; since most shadow cells are legitimately empty, this inflates both model and baselines, non-uniformly across sub-masks. We therefore report **micro-averaged** shadow IoU — pool integer intersection and union counts across all frames and divide once; empty denominators are undefined, never 1.0 — and the defended quantity is the **gain** over a persistence baseline (the observed frame carried forward), not the absolute level.

Completion vs. forecasting decomposition. All horizons are rendered from the same issue pose (static geometry projects to identical cells at every horizon), so each shadow cell is classified by its occupancy pattern across horizons: **STATIC** — occupied at *all* horizons (persistent hidden structure; $\approx 7\%$ of shadow), where recall measures **amodal completion**; **DYNAMIC** — occupied at *some but not all* horizons ($\leq 1.5\%$), where IoU measures genuine **forecasting** beyond persistence; **NEVER** —

tier (fixed val, @5 s)	gain	95% CI
static	+0.306	[+0.288, +0.327]
moving	+0.363	[+0.349, +0.378]
debris	+0.315	[+0.281, +0.350]




Figure 2: Micro shadow-IoU gain over persistence with 95% per-episode bootstrap CIs: per tier at 5 s, fixed geometry (left) and across horizons (right); every interval is clear of zero.

occupied at *no* horizon ($\approx 92\%$), the false-positive test. This separates completion from clairvoyance — a distinction pooled shadow-IoU hides.

Per-episode bootstrap CIs. Frames within an episode share geometry and pose, so we resample episodes with replacement ($B = 10,000$) and report the 2.5/97.5 percentile interval of the paired model-minus-persistence pooled micro-gain difference across replicates for every headline gain.

Calibration in shadow. We measure expected calibration error (ECE) restricted to shadow cells, contrasted with observed cells, using equal-mass binning (a fine 1000-bin histogram merged into 15 equal-count macro-bins, since in-shadow probabilities cluster near the prior and would leave equal-width bins empty).

4 Experiments

Setup. A MuJoCo single-room benchmark with three tiers — **static** (fixed obstacles), **moving** (kinematic movers), **debris** (objects that fall and settle) — 400 training / 80 validation episodes on disjoint seeds; per-tier analyses use 30-episode (9,000-frame) single-tier sets, randomized-geometry validation 80 episodes, and fixed-geometry results a 6×6 m room unless noted. Reproducibility is pinned: scene generation is bit-exact against a golden hash, and a rollout re-run reproduces every accuracy metric bit-identically (hardware-dependent latency excepted).

Shadow-gain headline. The model beats persistence at every horizon (Fig. 2), with every bootstrap interval clear of zero — statistically significant, not sampling noise. In policy rollouts the frame-averaged (macro) shadow gain is $+0.325/+0.322/+0.320/+0.178$ at 1/3/5/10 s. (Raw macro shadow-IoU levels — model 0.744 vs persistence 0.424 at 5 s — are empty-credit inflated and reported only for reference; the micro-gain is the defended number.)

Observation-blind prior. Persistence could be a weak baseline for *completion*, so we also evaluate a pose-marginal dataset prior (mean training-set occupancy in the egocentric frame, no observation at all), swept over thresholds and scored at its most favorable one. It reaches only 0.08–0.10 micro shadow IoU across tiers (0.09 on the randomized-geometry set): to recall hidden structure (0.87–0.92) it must blanket 72–79% of genuinely empty hidden space with false positives, whereas the model achieves 0.61–0.72 IoU at a 1.4–3.4% false-positive rate. Completion is thus observation-conditional inference, not a memorized average room; with geometry randomization (below), this rules out memorized layouts and marginals alike.

Completion, not clairvoyance. Decomposing the shadow mask attributes the gain almost entirely to amodal completion (Table 1). Excluding all persistent structure, the residual “nostatic” gain sits at the noise floor: fixed static +0.000, moving +0.007 (CI [+0.002, +0.012]), randomized-geometry static +0.002 ([+0.001, +0.004]) at 5 s. The honest claim is **amodal completion of hidden structure**, plus a small but statistically significant moving-tier forecasting signal that decays with horizon as genuine forecasting should. We do not claim dynamic forecasting on debris, where the dynamic signal is at the noise floor and persistence matches the model beyond 1 s.

Robustness to room geometry (memorization confound-kill). A fixed room shell could let the model memorize a constant wall prior, so we regenerated data with per-episode random width and depth (i.i.d. in [5, 8] m, obstacle counts scaled to keep $\sim 7\%$ positives), retrained from scratch, and evaluated on 80 held-out randomized rooms (Table 1, bottom): the gain does not drop (intervals

Table 1: Top: shadow-mask decomposition at 5 s, fixed geometry (dynamic IoU spans 1 s→10 s): the gain is recall of persistent structure, not forecasting. Bottom: static-tier gain, fixed vs per-episode randomized geometry.

component (fixed val @5 s)	model	persistence
persistent-structure recall (amodal completion)	0.83–0.93	0.37–0.51
dynamic IoU, moving tier (1 s→10 s)	0.092 → 0.031	0.060 → 0.020
dynamic IoU, debris tier	0.020 → 0.019	0.006 → 0.023
false-positive rate on never-occupied cells	0.014–0.034	0.022–0.027

static tier, @5 s	micro shadow gain	95% CI	completion recall
fixed 6 × 6 geometry	+0.306	[+0.288, +0.327]	0.829
randomized geometry	+0.309	[+0.294, +0.324]	0.854

overlap almost entirely) and completion recall is if anything higher on never-seen rooms: completion is per-scene inference, not a memorized shell.

Calibration in shadow. At 5 s, in-shadow ECE is 0.043–0.047 across tiers — *better* than in observed space (0.058–0.086) — with Brier 0.022–0.036; excluding persistent structure entirely, in-shadow ECE remains 0.055–0.056 on the tiers with dynamic content (the static tier’s structure-excluded region has no positive labels, hence no curve), so the calibration is not carried by the easy always-occupied cells. The model is underconfident in the mid-range but places little mass there; calibration is no fixed-room artifact: on randomized geometry, in-shadow ECE is 0.050–0.067 vs observed 0.056–0.077, with shadow the better-calibrated region at 5–10 s.

Ablations. Retraining without either auxiliary channel leaves the gain intact: no-visibility per-tier gains at 5 s (+0.309/+0.360/+0.313 vs full +0.306/+0.363/+0.315) are within CI and the no-flow rollout gain (+0.381 vs +0.320) is if anything higher; both are redundant (no optical flow on the critical path).

System. The perception-to-planning path runs at 5.97 ms p50 / 7.00 ms p95 (batch 1, deterministic U-Net), inside a 50 ms (20 Hz) budget; the excluded monocular-depth stage is not free (Depth-Anything-V2-Small: 49.9 ms p50 / 50.7 ms p95 in isolation, 54.4 ms measured end-to-end sequential, evaluation-era GTX 1080 Ti, 480×640), so deployment runs depth asynchronously while the 7 ms completion path consumes the latest depth frame at 20 Hz (audio synthesis excluded). A falling-object *proxy* (cell-level detection 0.787 at 2.62 s mean lead; component-level 0.752, precision only 0.714) counts events per due prediction, not per object, keys on cells revealed rather than newly occupied (occluded static geometry revealed by agent motion also counts), and its lead is bounded by the {1,3,5,10} s horizon grid, missing the ≥3 s target; we report it as a diagnostic, not a strength. Navigation: per-episode collision rate 0.367 (vs 0.60 untrained; straightness 0.773), missing the ≤0.10 target — an honest demonstration, no more.

Generative variant. A conditional-diffusion (DDPM) world model is ongoing work: iterated DDIM sampling exceeds our batch-eval budget; a like-for-like comparison is future work.

5 Related work

Predicting the structure of unobserved space has been approached from three angles. Semantic scene completion, from SSCNet [Song et al., 2017] through MonoScene [Cao and de Charette, 2022] and its monocular successors VoxFormer [Li et al., 2023] and VisHall3D [Lu et al., 2025], infers dense voxel geometry — and even hallucinates geometry beyond the visible frontier — from a single image, but targets outdoor driving voxels evaluated purely by mIoU/IoU, without calibrated uncertainty in the occluded region. Occupancy world models such as OccWorld [Zheng et al., 2024], together with occupancy-forecasting benchmarks (Occ3D; Tian et al., 2023) and self-supervised raycasting forecasters [Khurana et al., 2023], instead predict the *temporal* evolution of multi-frame occupancy for autonomous vehicles. Closest in spirit, ProxMaP [Sharma et al., 2023] completes indoor proximal occupancy from a single view for navigation, yet emits only point estimates. ShadowWeave is distinct:

147 egocentric BEV amodal completion into explicitly masked occluded space from one monocular-depth
148 frame, with calibrated uncertainty inside the hidden region.

149 6 Conclusion and limitations

150 Calibrated amodal completion of occluded occupancy works from one monocular-depth frame,
151 transfers to never-seen rooms, and fits a real-time perception-to-planning budget; the empty-credit-
152 immune micro-gain, per-episode CIs, and completion/forecasting decomposition keep the claim
153 honest and, we hope, get reused. The limitations are those it surfaces: evaluation is in simulation
154 (ground-truth occluded occupancy is unavailable on real footage without instrumentation), so we
155 make no quantitative real-world claim, though the front-end already accepts real monocular depth
156 (Depth-Anything-V2), making sim-to-real transfer the natural next step; the model performs amodal
157 *completion*, not multi-step *forecasting*, and is framed accordingly; navigation misses its collision
158 target and the audio interface is not user-evaluated (system context); and falling-object lead is bounded
159 by the horizon grid.

160 References

- 161 Anh-Quan Cao and Raoul de Charette. Monoscene: Monocular 3d semantic scene completion. In
162 *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*,
163 2022. arXiv:2112.00726.
- 164 Tarasha Khurana, Peiyun Hu, David Held, and Deva Ramanan. Point cloud forecasting as a proxy for
165 4d occupancy forecasting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and*
166 *Pattern Recognition (CVPR)*, 2023. arXiv:2302.13130.
- 167 Yiming Li, Zhiding Yu, Christopher Choy, Chaowei Xiao, Jose M. Alvarez, Sanja Fidler, Chen Feng,
168 and Anima Anandkumar. Voxformer: Sparse voxel transformer for camera-based 3d semantic
169 scene completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
170 *Recognition (CVPR)*, 2023. arXiv:2302.12251.
- 171 Haoang Lu, Yuanqi Su, Xiaoning Zhang, Longjun Gao, Yu Xue, and Le Wang. Vishall3d: Monocular
172 semantic scene completion from reconstructing the visible regions to hallucinating the invisible
173 regions. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*,
174 2025. arXiv:2507.19188.
- 175 Vishnu Dutt Sharma, Jingxi Chen, and Pratap Tokekar. Proxmap: Proximal occupancy map prediction
176 for efficient indoor robot navigation. In *2023 IEEE/RSJ International Conference on Intelligent*
177 *Robots and Systems (IROS)*, 2023. arXiv:2305.05519.
- 178 Shuran Song, Fisher Yu, Andy Zeng, Angel X. Chang, Manolis Savva, and Thomas Funkhouser.
179 Semantic scene completion from a single depth image. In *Proceedings of the IEEE Conference on*
180 *Computer Vision and Pattern Recognition (CVPR)*, 2017. arXiv:1611.08974.
- 181 Xiaoyu Tian, Tao Jiang, Longfei Yun, Yucheng Mao, Huitong Yang, Yue Wang, Yilun Wang, and
182 Hang Zhao. Occ3d: A large-scale 3d occupancy prediction benchmark for autonomous driving. In
183 *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*,
184 2023. arXiv:2304.14365.
- 185 Wenzhao Zheng, Weiliang Chen, Yuanhui Huang, Borui Zhang, Yueqi Duan, and Jiwen Lu. Occworld:
186 Learning a 3d occupancy world model for autonomous driving. In *European Conference on*
187 *Computer Vision (ECCV)*, 2024. arXiv:2311.16038.